

ClipMind AI — Presentation Narration Script (10 Slides)

Slide-by-Slide Verbal Explanation & Presentation Guide

Project Title: ClipMind AI: Video Summarization & Key Moments Detection Platform **Presenter:** Adabala Venkata Thrinadh **Course & Institution:** B.Tech in Computer Science & Engineering (Batch 2024–2028), GMR Institute of Technology, Rajam **Target Duration:** 8 to 10 Minutes **Deck Design:** 10 High-Impact, Diagrammatic, and Flowchart-Rich Slides with Minimal Text

Slide 1: Title & Executive Overview

Verbal Narration:

""Respected evaluators, faculty mentors, and fellow engineers, good morning/afternoon. My name is **Adabala Venkata Thrinadh**, a B.Tech Computer Science and Engineering student from **GMR Institute of Technology, Rajam**. Today, I am honored to present **ClipMind AI: A Video Summarization and Key Moments Detection Platform**, engineered under the **Infosys Springboard AI & Full-Stack Cloud Internship**. In today's digital learning environment, video accounts for over 80% of web traffic. However, video remains an inherently linear, continuous medium. Our mission with ClipMind AI is to transform passive video recordings into structured, searchable, and interactive knowledge assets with persistent zero-loss cloud storage and interactive concept mind maps.""

Slide 2: Problem Statement & Industry Motivation

Verbal Narration:

""On Slide 2, we highlight the four critical industry bottlenecks in video learning: - **Linear Scrubbing Friction:** Manually scrubbing through a 60-minute lecture to locate a 2-minute formula wastes up to 80% of study time. - **Absence of In-Video Semantic Search:** Standard media players cannot search through spoken dialogue, mathematical definitions, or visual slide transitions. - **Ephemeral Cloud Media Loss:** Cloud containers on hosts like Render or AWS ECS wipe local disks during restarts, causing uploaded media to be lost. - **Passive Cognitive Decay:** Passive video consumption leads to rapid memory degradation without active recall and conceptual synthesis. Opposite these bottlenecks, ClipMind AI introduces sub-second word-level seeking, dual-tier summarization, Google Drive zero-loss cloud storage with sub-20ms streaming, and active recall learning.""

Slide 3: Proposed Solution & Novelty (6 Architectural Pillars)

Verbal Narration:

""Slide 3 presents our architectural novelty matrix—six foundational engineering pillars:* 1. **Google Drive 15GB Cloud Storage & HTTP 206 Streaming:** Direct OAuth 2.0 GIS token exchange gives every user 15 GB of persistent personal cloud storage with sub-20ms byte-range seeking. 2. **Universal Pipeline Cloud Fallback:** If local container disk is reset, our pipeline automatically downloads media from Google Drive into a high-speed working buffer for seamless 7-stage processing. 3. **Interactive AI Concept Mind Maps:** Synthesizes lecture transcripts into a hierarchical visual knowledge graph with click-to-seek video synchronization. 4. **Multimodal Video Intelligence Engine:** Integrates OpenAI Whisper ASR (4.18% WER) with OpenCV 1 fps visual frame differencing for slide cut detection. 5. **Dual-Tier NLP Summarization:** Combines LexRank graph centrality (<1.5s) for instant TL;DRs with generative LLMs for modular chapters and quizzes. 6. **5-Tab Educator Studio & Active Recall Suite:** An authoring workspace providing transcript editing, chapter structuring, quiz creation, flashcards, and student preview.""

Slide 4: Architecture Structural Flow & Layer Interconnections

Verbal Narration:

""Slide 4 diagrams our five-layer decoupled, service-oriented architecture:* - **Layer 1: Frontend Client Presentation:** Built with React 19, TypeScript 5.5, and Vite 8.2, featuring a bespoke engineering UI and SVG iconography. - **Layer 2: API Gateway & Security:** FastAPI ASGI server coordinating OAuth2 JWT authentication, Google Identity Services, declarative RBAC guards, and HTTP 206 streaming. - **Layer 3: Asynchronous Media & AI Engine:** FFmpeg demuxes audio to 16kHz mono WAV, Whisper transcribes with word timestamps, LexRank/LLMs summarize, and OpenCV detects scene cuts. - **Layer 4: Polyglot Persistence:** SQLite/PostgreSQL provides ACID compliance for user authentication, while MongoDB Atlas stores unstructured intelligence documents. - **Layer 5: Zero-Loss Cloud Storage:** Interfaces with Google Drive API v3 and high-speed temporary buffers to guarantee zero data loss.""

Slide 5: End-to-End Processing Workflow (Complete Data Lifecycle)**Verbal Narration:**

""Slide 5 traces the complete end-to-end data lifecycle across five operational phases:* - **Phase 1: Ingestion Sources:** Supports local MP4, MOV, MKV, WebM uploads, YouTube URL ingestion via yt-dlp, or Google Drive cloud media. - **Phase 2: Gateway & Cloud Storage:** Validates MIME types, assigns UUIDs, triggers background auto-backup to Google Drive, and opens WebSocket connections. - **Phase 3: Multimodal AI Inference:** Sequentially extracts 16kHz WAV audio, transcribes via Whisper, extracts LexRank and LLM chapters, and calculates OpenCV HSV histograms. - **Phase 4: Persistence & Delivery:** Writes structured intelligence to MongoDB Atlas, marks video completed, and configures the HTTP 206 byte-range proxy. - **Phase 5: Role Workspaces:** Dispatches tailored views to Content Creators, Learners, Educators, and Administrators.""

Slide 6: Detailed 7-Stage Video Processing Pipeline**Verbal Narration:**

""Slide 6 provides granular detail on our 7-stage processing pipeline with sub-second WebSocket telemetry:* - **Stage 1 (15% - Ingestion & Probe):** Validates MIME types and stages media in working buffers. - **Stage 2 (35% - Media Demuxing):** FFmpeg converts video audio to 16kHz mono 16-bit PCM WAV. - **Stage 3 (65% - Whisper Speech-to-Text):** Generates word-level timestamps across 80-channel Log-Mel spectrograms. - **Stage 4 (75% - Dual-Tier Summarization):** LexRank graph centrality produces instant TL;DR while LLMs synthesize chapters. - **Stage 5 (85% - OpenCV Vision Cuts):** Detects slide transitions and visual shifts at 1 fps, saving WebP thumbnails. - **Stage 6 (90% - Concept Mind Map):** Structures lecture concepts into a hierarchical knowledge tree. - **Stage 7 (100% - Persistence & HTTP 206):** Commits documents to MongoDB Atlas and activates sub-20ms seeking playback.""

Slide 7: Milestone-Wise Technical Implementation Roadmap (Weeks 1–8)**Verbal Narration:**

""Slide 7 showcases our 8-week engineering lifecycle, executed with 100% adherence to the Infosys Springboard Specification:* - **Milestone 1 (Weeks 1 & 2):** Architecture, Polyglot Database, Salted Bcrypt Auth, and Chunked Video Ingestion. - **Milestone 2 (Weeks 3 & 4):** OpenAI Whisper STT, Sub-Second Word Timestamps, LexRank Extractive Summarizer, and LLM Chapters. - **Milestone 3 (Weeks 5 & 6):** OpenCV Scene Cuts, 5-Tab Educator Studio, Learner Study Room, Multi-Format Exports, and WebSocket Telemetry. - **Milestone 4 (Weeks 7 & 8):** Google Drive 15GB Cloud Integration, HTTP 206 Range Proxy, Interactive Concept Mind Maps, Docker Containerization, and 100% Test Pass Rate.""

Slide 8: Multi-Persona Role-Based Access Control (RBAC) & Educator Studio

Verbal Narration:

""Slide 8 highlights our four dedicated stakeholder personas and our Educator Studio: - **Content Creators:** Manage video libraries, monitor AI processing, view analytics, and export packages in PDF, DOCX, TXT, SRT, and VTT. - **Learners:** Benefit from synchronized playback with auto-scrolling transcripts, click-to-see, auto-graded quizzes, and 3D flip flashcards. - **Educators:** Utilize our **5-Tab Authoring Studio** to: - Tab 1: Edit transcripts and adjust speaker diarization tags. - Tab 2: Organize modular curriculum chapters and time bounds. - Tab 3: Author interactive quizzes with answer keys and explanations. - Tab 4: Create active recall flashcard decks. - Tab 5:* Test in a WYSIWYG Student Preview mode before publishing. - **Administrators:** Oversee user governance, role privileges, 50+ event audit logs, and one-click cache maintenance.""

Slide 9: Quantitative Benchmarks, Telemetry & Performance Evaluation

Verbal Narration:

""Slide 9 presents our verified empirical benchmarks:* - **4.18% Word Error Rate (WER)** with Whisper STT (95.82% speech recognition accuracy). - **46.8% ROUGE-1 and 42.1% ROUGE-L** summarization scores capturing human-level core topics. - **92.4% Precision** in OpenCV visual presentation slide transition detection. - **2.6× Real-Time Speedup**, transcribing and analyzing an hour of lecture video in just 2.6 minutes. - **< 20ms Seek Latency** via our custom HTTP 206 Partial Content video streaming proxy. - **100% Automated Test Pass Rate** across all 28/28 end-to-end integration tests.""

Slide 10: Conclusion, Technical Stack & Operational Readiness

Verbal Narration:

""To conclude on Slide 10, ClipMind AI delivers an enterprise-grade, production-ready solution that transforms video learning: - Fully containerized with multi-stage Dockerfiles and docker-compose. - Eliminates video data loss permanently through Google Drive cloud integration. - Combines acoustic speech recognition, natural language processing, computer vision, and knowledge graphs into a unified reactive platform. Thank you very much for your time and guidance. I am now ready to present our live application walkthrough and answer your questions.""